
ADL x MLDS 2017 Fall

HW3 - Game Playing

2017/11/19
adlxmls@gmail.com

Outline

- **Introduction**
 - Game Playing
 - Game Type
- **Deep Reinforcement Learning**
 - Policy Gradient
 - Deep Q-Learning (DQN)
 - Improvements to Policy Gradient & DQN
- **Grading & Format**
 - Grading Policy
 - Code Format
 - Report
 - Submission

Introduction


Game Playing

- Implement an agent to play Atari games using Deep Reinforcement Learning
- In this homework, you are supposed to implement **Policy Gradient** and **Deep Q-Learning** (DQN)

Introduction

Environment

Pong



Breakout




<https://gym.openai.com/envs/>

Deep Reinforcement Learning

Policy Gradient

REINFORCE algorithm:

- 
1. sample $\{\tau^i\}$ from $\pi_\theta(\mathbf{a}_t|\mathbf{s}_t)$ (run it on the robot)
 2. $\nabla_\theta J(\theta) \approx \sum_i (\sum_t \nabla_\theta \log \pi_\theta(\mathbf{a}_t^i|\mathbf{s}_t^i)) (\sum_t r(\mathbf{s}_t^i, \mathbf{a}_t^i))$
 3. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$

Deep Reinforcement Learning


REINFORCE Baseline

1. Training loop(simplest version):
 - a. Play until a game is over(one player gets 21 points) with policy network π_θ and store (s,a,r) tuples into memory m.
 - b. Discount and normalize rewards in memory into r' to reduce variance
 - c. Approximate gradient $\nabla_\theta J(\theta) \approx \sum_{(s_t, a_t, r'_t) \in m} \nabla_\theta \log \pi_\theta(a_t | s_t) r'_t$
 - d. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
 - e. Clear the memory m
2. Tips:
 - a. The trajectory length varies from game to game, hence sum the gradient instead of averaging it.
 - b. Feed $s'_t = s_t - s_{t-1}$ into policy network, where s_t comes from environment at time step t and $s'_0 = s_0$
 - c. When one player gets point, reset the running add of discounted reward to zero

Deep Reinforcement Learning

Deep Q-Learning (DQN)

“classic” deep Q-learning algorithm:

- 
1. take some action \mathbf{a}_i and observe $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$, add it to \mathcal{B}
 2. sample mini-batch $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$ from \mathcal{B} uniformly
 3. compute $y_j = r_j + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$ using *target* network $Q_{\phi'}$
 4. $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_\phi}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_\phi(\mathbf{s}_j, \mathbf{a}_j) - y_j)$
 5. update ϕ' : copy ϕ every N steps

Deep Reinforcement Learning

Improvements to Policy Gradient (BONUS)

- Variance Reduction
- Advanced Advantage Estimation
- Off-policy learning by Importance Sampling
- Natural Policy Gradient
- Trust Region Policy Optimization
- Proximal Policy Optimization

http://rll.berkeley.edu/deeprlcourse/f17docs/lecture_4_policy_gradient.pdf

http://rll.berkeley.edu/deeprlcourse/f17docs/lecture_13_advanced_pg.pdf

Deep Reinforcement Learning

Improvements to DQN (BONUS)

- Double Q-Learning
- Dueling Network
- Prioritized Replay Memory
- Multi-Step Learning
- Noisy DQN
- Distributional DQN

<https://arxiv.org/pdf/1710.02298.pdf>

Grading Policy

- Baseline (6%)
 - Policy Gradient (3%)
 - Getting averaging reward in 30 episodes over **1** in **Pong**
 - DQN (3%)
 - Getting averaging reward in 100 episodes over **40** in **Breakout**
- Report (10%)
- Bonus (4%)

Grading & Format

Baseline (6%)

- Policy Gradient (3%)
 - Getting averaging reward in 30 episodes over **1** in **Pong**
 - Without OpenAI's Atari wrapper & reward clipping
 - Improvements to Policy Gradient are allowed, including Actor-Critic series
- DQN (3%)
 - Getting averaging reward in 100 episodes over **40** in **Breakout**
 - With OpenAI's Atari wrapper & reward clipping

Grading & Format

Code Format

- Please download the sample files from [github](#)
- Follow the instructions in README to install required packages
- **Four** functions you should implement in [agent_\[pg|dqn\].py](#)
 1. `__init__(self, env, args)`
 2. `init_game_setting(self)`
 3. `train(self)`
 4. `make_action(self, state, test)`
- **DO NOT** add any parameter in `__init__()`, `init_game_setting()` and `make_action()`
- You can add new methods in the [agent_\[pg|dqn\].py](#)
- You can add your arguments in [argument.py](#)

Grading & Format

Report (10%)

- Up to 6 pages, in Chinese
- Basic Performance (6%)
 - Describe your Policy Gradient & DQN model (1% + 1%)
 - Plot the learning curve to show the performance of your Policy Gradient on Pong (2%)
 - Plot the learning curve to show the performance of your DQN on Breakout (2%)
 - X-axis: number of time steps
 - Y-axis: average reward in last 30 episodes

Grading & Format

Report (10%)

- Experimenting with DQN hyperparameters (4%)
 - Choose one hyperparameter of your choice and run at least three other settings of this hyperparameter
 - You should find a hyperparameter that makes a nontrivial difference on performance
 - Plot all four learning curves in the same graph (2%)
 - Explain why you choose this hyperparameter and how it effect the results (2%)
 - Candidates: learning rate, gamma, network architecture, exploration schedule/rule, target network update frequency, etc.

Grading & Format

Bonus (4%)

- You can train on any environment to show your results
- Improvements to Policy Gradient (2%)
 - Implement at least **two** improvements to Policy Gradient (p.8) and describe why they can improve the performance (1%)
 - Plot a graph to compare and analyze the results with and without the improvements (1%)
- Improvements to DQN (2%)
 - Implement at least **two** improvements to DQN (p.9) and describe why they can improve the performance (1%)
 - Plot a graph to compare and analyze the results with and without the improvements (1%)
- Implement other advanced RL method, describe what it is and why it is better (2%)
 - Ex: Actor-Critic, A2C, A3C, ACKTR
- Up to 4 bonus points

Grading & Format

Late submission

- Please fill the [late submission form](#) first **only if you will submit HW late**
- Please push your code before you fill the form
- **There will be 25% penalty per day for late submission,** so you get 0% after four days
- You get 0% if the required files has bug.
 - If the error is due to the format issue, please come to fix the bug at the announced time, or you will get 10% penalty afterwards.

Grading & Format

Submission

- Deadline: **2017/12/16 23:59 (GMT+8)**
- Your github **MUST** have 5 files under directory hw3/
 - agent_dir/agent_pg.py
 - agent_dir/agent_dqn.py
 - [saved_model_file] * 2
 - report.pdf
 - argument.py (optional)
 - README (optional)
 - download.sh (optional)
 - other files you need
- If your model is too large for github, upload it to a cloud space and write download.sh to download the model
- Do not upload any file named the same with other sample codes

Grading & Format

Grading

- Please use Python with version ≥ 3.5
- The TAs will execute `'python3 test.py --test_pg --test_dqn'` to run your code
- The execution for both model should be done within 10 minutes respectively, excluding model download
- Allowed packages:
 - PyTorch v0.2.0
 - Tensorflow r1.3
 - Keras 2.0.7 (Tensorflow backend only)
 - MXNet 0.11.0
 - CNTK 2.2
 - Numpy
 - Pandas
 - Python Standard Lib
- If you use other packages, please ask for permission first !!!

Related Materials

- Course & Tutorial:
 - [Berkeley Deep Reinforcement Learning, Fall 2017](#)
 - [David Silver RL course](#)
 - [Nips 2016 RL tutorial](#)
- Blog:
 - [Andrej Karpathy's blog](#)
 - [Arthur Juliani's Blog](#)
- Text Book:
 - [Reinforcement Learning: An Introduction](#)

TA Information

TA hours

- 有問題請利用TA hours、信箱或FB社團，**請不要FB私訊助教！！**
- If you have other questions,
 - please contact TAs via adlxmlds@gmail.com
 - post your questions on [facebook group](#)
 - go to TA office hours
 - 陳璽安 Thur 14:00-15:30 (德田536)
 - 王耀賢 Fri 16:00-17:30 (電二531) (11/30開始)
 - 葉奕廷 Mon 10:30-12:00 (德田524) (12/4請假)